Contents

0.1 Grace-Mar — Business Prospectus

Identity Infrastructure for the Agent Web

Version 1.0 · February 2026

0.1.1 Executive Summary

Grace-Mar provides the **identity substrate for the agent web**: a structured, evidence-grounded Record of who a person is, with a user-controlled gate that ensures only approved content enters. As agents become economic actors—searching, paying, executing across services—they need trusted identity data. Today, that primitive does not exist at scale.

Positioning: Grace-Mar is not a second-brain tool, AI clone, or child-specific tutor. It is **identity infrastructure**—the trust primitive agents and AI schools need to know who they serve.

Dual opportunity: 1. **Supplemental** — Add user-owned Record layer to Alpha ($40K–$75K/year AI schools) and similar platforms. Record feeds personalization; school events feed pipeline. 2. **Low-cost alternative** — Open-source. Families outside elite tuition run Grace-Mar + Khan/IXL + lightweight structure for $0 software cost.

0.1.2 Problem

Pain Point	Impact
Agent web has no identity primitive	Agents need structured, trustworthy identity; today they scrape, infer, or hallucinate
AI schools own the data	Alpha, Khan, IXL store student data in proprietary systems; families have no portable asset
Evidence ≠ artifacts	Platforms use engagement metrics; nobody curates artifacts (writing, art, "we did X") with user approval
$40K+ tuition excludes most families	Elite AI schools prove demand; 99% of families cannot afford them

---

0.1.3 Solution

Record — Structured, evidence-grounded identity: who they are (SELF) + what they can do (SKILLS).

Gate — User approves every merge. The agent may stage; it may not merge. Trust boundary is architectural, not configurable.

Evidence grounding — Every claim traces to artifacts (ACT-XXXX). No LLM inference; no undocumented knowledge.

Export — Agent-native interface. Markdown, schema, manifest. Identity can be consumed by Alpha/Incept, OpenClaw, and future agent ecosystems.

0.1.4 Market Opportunity

Segment	Size	Grace-Mar Fit
Elite AI schools	Alpha $40K–75K; Astra Nova up to $75K	Supplemental Record layer; integration license

Segment	Size	Grace-Mar Fit
Mid-tier microschools	Prisma $5K–12K; Synthesis ~$2K/year	Supplemental or alternative
Khan/IXL families	Khan free; Khanmigo $4/mo; IXL $20/mo	Low-cost alternative; open-source core
Homeschool	Growing; Alpha homeschool = 1x (motivation bottleneck)	Record + "we did X" = motivation substrate

Gap: No platform provides user-owned, evidence-grounded, portable identity. Grace-Mar fills this across every segment.

0.1.5 Traction & Status

Milestone	Status
Architecture	Complete — Record schema, gated pipeline, three-channel mind (Knowledge, Curiosity, Personality)
Pilot	Active — pilot-001 (age 6) post-seed; SELF, SKILLS, EVIDENCE populated
Emulation	Live — Telegram bot; WeChat optional
Pipeline	Active — analyst stages candidates; user approves; merge to profile
Export	`export_user_identity.py` → USER.md/SOUL.md for OpenClaw

0.1.6 Monetization Paths

Path	Value	Revenue Model
Supplemental (Alpha families)	Record layer for AI schools	Integration license; subscription for Record hosting + export
Low-cost alternative	Open-source for families outside elite tuition	Open core free; hosted service; support; premium export
B2B (AI school platforms)	Identity substrate for Alpha, future schools	Platform license; API fees; white-label
Infrastructure	Open schema, protocol, trust primitives	Reference implementation; certification; ecosystem

---

0.1.7 Competitive Moat

Moat	Description
Evidence grounding	Every claim traces to artifacts. Competitors use engagement metrics.
Gate as trust boundary	User sovereignty architectural. Schools auto-update; Grace-Mar never does.
Agent-native interface	Record structured for consumption; export, schema, manifest.
Portable architecture	Principles and schema adoptable; Grace-Mar = reference implementation.
Open-source + low-cost	No $40K tuition. Homeschoolers, public-school families can adopt.

0.1.8 Ask

Partners: AI school platforms (Alpha, Prisma, Synthesis) seeking identity layer; OpenClaw and agent ecosystems needing identity source.

Investors: Pre-seed / seed. Use of funds: hosted service, homeschool documentation, Alpha/Incept integration path, team.

0.1.9 Risks & Mitigations

Risk	Mitigation
Agent web adoption slower than expected	Grace-Mar works standalone (Record + Voice); identity value persists
Competitor copies schema	First-mover; evidence depth; governance maturity; open protocol grows ecosystem
Trust incidents	Document security posture; treat agent as adversary; gate non-negotiable
User fatigue (approval burden)	Staging automation; lightweight approve/reject; no editing required

---

0.1.10 Acknowledgments

The ideas behind Grace-Mar draw on the work of: Alexander Wissner-Gross (causal entropic forces), Peter Diamandis (abundance), Nick Bostrom (superintelligence), Ray Kurzweil (singularity), Brian Roemmele (multimodal AI), Scott Adams (systems thinking), Julian Jaynes (bicameral mind), and Satoshi Nakamoto (decentralized trust).

0.1.11 Key Documents

Document	Purpose
White Paper	Full narrative, differentiation, technical model
Design Notes	Positioning, agent-web insights, proposal angles
Business Roadmap	Strategy, priority roadmap, metrics
Architecture	Technical design, pipeline, modules

Grace-Mar · A cognitive fork — versioned, evidence-grounded, user-owned